

When Does a Consequence Model Make AI Agents Safer?

Pre-registered Studies on Demo Apps, Real Self-Hosted Apps and a Real Terminal

Caio Vicentino
OpenInterpretability
Fortaleza, Brazil
ORCID: 0009-0003-4331-6259
caio@openinterp.org

October 2026

DOI: [10.5281/zenodo.23197341](https://doi.org/10.5281/zenodo.23197341)

Abstract

An agent that cannot see what an action will change will sometimes do the wrong thing. We give agents a tool that asks Ekibasis-27B, an open consequence model that answers in one forward pass, what an action would change, and measure harm and task success in studies whose plans were hashed before the runs (locally; the documents are released). On real, unmodified Gitea, Nextcloud and Roundcube driven through a real browser, Claude Sonnet 5.5’s harmful runs fell from 25.0% to 4.2% in a pre-registered secondary contrast with one run per task (−20.8 points, 95% CI over tasks $[-37.5, -6.2]$, over the 9 task families $[-47.8, 0.0]$), mostly on tasks whose request names the harmful action, with more runs stopping short. The pre-registered hypotheses for Claude Haiku 4.5 failed: warned, it mostly did what the request said. In a pre-registered follow-up on 32 fresh tasks, a guard that offers a safer way Ekibasis has checked and otherwise asks the user cut Haiku’s harmful runs from 72.9% to 2.1% with a scripted user (Sonnet’s, in a secondary arm, from 50.0% to 8.3%). Asking before every consequential click without Ekibasis did as well (the scripts had been checked to object to its questions on each harmful path), and the guard asked 43% as many questions where the plan’s bar was a quarter, so that hypothesis failed. On our demo apps, harmful runs fell for all five models we tested (Sonnet: 24/60 to 0/60) and a placebo did not help Sonnet; Haiku and Qwen 9B alone had done about as well as Sonnet, so the gains show the forecast’s effect, not a weaker agent lifted above a stronger one (GLM’s runs alone are not comparable: see the erratum). Three further pre-registered studies failed their bars, where the agent alone was at the ceiling or the policy stated the consequences. In a real terminal, a Claude Code hook asked on 2.9 per 100 commands of a careful agent and kept Haiku’s planted work in 9/9 sessions against 5/9 and 7/9 without it in two small studies (one exploratory). In the one-pass format Ekibasis was trained on, it was more accurate than Qwen-AgentWorld-35B-A3B on all 12 of our suites; allowed to reason, AgentWorld won 1, lost 2 where its reasoning hit the token cap, and showed no resolved difference on 9. The tasks are ours and are released with the app setups, adapters, judges, tables and scripts.

Erratum (v2, 10 October 2026). In the GLM-5.3-Flash and Qwen 3.5 runs, the agent harness (opencode) also loaded a `computer-use` skill from the operator’s user directory. It was not part of the study, and it tells the agent never to move money. In five of GLM’s six bank runs without Ekibasis the agent loaded it and stopped before paying. GLM’s low harm alone and its stopped-short runs came from this skill, so v1’s claims that GLM alone was the most cautious agent and did less harm than Sonnet alone are withdrawn. Every harm difference with Ekibasis keeps its sign without

the affected runs, although for GLM and Qwen 4B the interval now reaches zero (Section 6.1). The Claude runs could not load skills. Changed in v2: the abstract, the introduction, Section 6, a new paragraph in Section 6.1, the limitations and the conclusion.

1 Introduction

A strong language model usually reasons well about what an action does when the screen shows it. It cannot reason about what the screen does not state: that a “Clean up” button also deletes real orders, that cancelling one flight of a round trip cancels both, that making a repository public exposes its history. The rules of the application know these things, and a model trained to apply such rules to a state can say what an action will change before the agent takes it. This paper asks whether such a consequence model, called before consequential actions, lowers the harm agents do without lowering what they get done; whether the gain holds across agents of different families; what it costs; and where it stops.

We use Ekbasis-27B [33], an open consequence model that reads the probability of every listed outcome of a question about the next state in one forward pass, and study it in five settings:

- **Hidden consequences on demo apps** (Section 5): 82 tasks in 12 small web apps written for this purpose, with placebo, Claude-as-forecaster and true-answer controls, and five models in six configurations (Section 6).
- **Real self-hosted apps** (Section 9): unmodified Gitea, Nextcloud and Roundcube with a local mail server, a real browser, and the truth read from the apps’ own interfaces.
- **Turning warnings into safe actions** (Section 10): fresh tasks on the same apps, with a guard that offers a safer way it has checked with Ekbasis and asks the user when there is none.
- **A real terminal** (Section 11): Claude Code on clones of real repositories with a hook that asks Ekbasis before git and shell commands.
- **Three studies whose bars failed** (Section 8): look-ahead on multi-screen apps and a state tracker, where the agent alone was at the ceiling, and τ -bench retail, whose policy states the consequences.

We also compare the consequence model itself with Qwen-AgentWorld-35B-A3B [44], an open world model trained on real agent trajectories (35B parameters, 3B active per token), on the same questions (Section 12).

The main results are these. With the consequence model in the loop, Claude Sonnet did less harm where the consequence was hidden: clearly on our demo apps, and on real apps in a secondary contrast carried by three task families. Haiku, GLM and Qwen 9B with it did less harm than Sonnet alone on our demo apps (Qwen 4B did not resolve), where alone Haiku and Qwen 9B did about as well as Sonnet (GLM’s runs alone are not comparable; see the erratum), so this shows the effect of the forecast, not a weaker agent lifted above a stronger one. On real apps Claude Haiku mostly did what the request said even when warned; a guard that offers a safer way checked by Ekbasis and asks the user when there is none cut its harmful runs on fresh tasks to 2.1%, with a simulated user (and Sonnet’s, in a secondary arm, from 50.0% to 8.3%), and asking the user every time, without Ekbasis, did as well with more questions; the simulated user’s scripts had been checked to object to that condition’s question on every harmful path. Three further studies failed their bars: two where the agent alone was at the ceiling, and τ -bench, whose policy already states the

consequences. The plans of these studies were written and hashed before their runs (the hashes were recorded locally; the documents are released), with these exceptions: the GLM and Qwen runs on the demo apps were not separately pre-registered, the look-ahead plan was not hashed, and the Haiku terminal sessions with client 0.1.2 and the Qwen arm on the real apps were registered as exploratory (Section 4). The per-run tables, the analysis script that recomputes every number in this paper from them and from three aggregated tables, and a checker that fails if the text cites a number the script does not produce (it checks that each number exists, not where it is used) are released with the paper.

2 Related work

World models for agents. Language-model agents have been given world models that predict the next observation, to look ahead before acting on the web [4, 15] or to train agents in simulation at scale [36, 44]. CWM trains a code model on execution traces to model program state [9]; agents can also learn from the consequences of their own early actions [42], and a world model can be trained so that the state it writes distinguishes the true outcome from those of alternative actions [21]. Earlier work asked whether language models can serve as text-based world simulators [35] and used one as a world model for planning [16]. Ekbasis does not write the next state: it scores the listed answers to a typed question about it in one pass, which makes its confidence available to a guard. Our earlier paper studied when an agent using it should look at the real state, and found that consequence training makes the model’s remaining errors confident [33]; we return to those errors in Section 6.2.

Side effects and pre-action guards. Several benchmarks measure the harm agents do or the rules they break. ToolEmu emulates tools with a language model to find the risks of agents whose user intent is benign [32]; OS-Harm, built on a computer-use environment, covers deliberate user misuse, prompt injection and model misbehavior [18]; ST-WebAgentBench scores web agents on enterprise tasks, each paired with safety and trustworthiness policies [20]; τ -bench [40] and τ^2 -bench [3] score whether tool-using agents follow a domain policy and reach the goal state, not harm as such; SABER measures the operational safety of coding agents in stateful workspaces [17], and a different SABER safeguards an agent’s mutating steps [8]; out-of-scope actions of coding agents on benign tasks have been measured directly [25], and a revisit of a workplace benchmark found unintended harmful actions, such as emailing the wrong person, to have become rare for frontier models [38]. Guards that act before the action include WebGuard [43], world-model-based guards for mobile GUI agents [41] and for LLM agents in general [11], and CORA, which bounds the rate of harmful actions with conformal risk control [7]. Simply telling an agent to quit when unsure also improves safety [27]; our placebo condition is a close relative. Our setting differs in two ways: the harm comes from a consequence the screen does not state, and the guard is a separate open model that returns a probability for each outcome and that any agent can call.

Selective prediction and guarantees. A guard that trusts its confident answers needs those answers to be right; Learn then Test [1] and conformal risk control [2] turn calibration data into guarantees, and KnowNo uses conformal prediction to make a planner ask for help when unsure [30]. This paper reports no such guarantee; it measures what agents do with the forecasts.

3 Setup

The consequence model. The released Ekbas-27B (a weight interpolation [39] of two fine-tunes, described in [33]): the weights of the Hugging Face revision `e0675dd` of 4 October 2026, unchanged since and listed with their SHA-256 in the repository’s `SHA256SUMS`; the studies’ plans record that the served copy was these released weights, and we did not hash the served copy again for this paper. It was served with vLLM [19] behind the released `serve.py` 1.3 and called through the released client’s prompts. A question gives the rules, the current state, the action and the listed outcomes; the model returns a probability for each outcome. It has no per-call charge; its GPU time is not priced here.

The foresee tool and the panel. In the demo-app studies the agent works on a desktop of small web apps through desk tools only (`look`, `open_app`, `click`, `type_text`, `say`, `done`) and, in every condition except *blind*, `foresee(element)`. For each consequential element the app supplies the panel’s specification: its rules in plain text, the current state, the action, and questions whose harmful answer is marked. The panel shows the predicted outcome, its confidence and a “harmful” mark when the predicted outcome is the marked one. Each demo app implements one mechanic in six task variants, and we wrote its questions, with the harmful answer marked, for that mechanic; the panel therefore tells the agent which outcome we consider harmful. On the real apps (Section 9) the questions are per action type and the state is read from each app’s interface at call time.

Conditions. *blind*: no tool. *see*: the tool answered by Ekbas. *placebo*: the same tool returns “Consider what this action will do before acting.” *oracle_llm*: Claude Sonnet answers the same panel questions. *oracle_truth*: the app’s own truth answers them. *router* (GLM and Qwen only): *see*, except that the tool server appends a note to use the agent’s own judgment to any panel with a confidence below 0.5. Every non-blind condition has the same tool description and system prompt; only what the tool returns differs. The real-apps studies add conditions described in Sections 9 and 10.

Tasks. On the demo apps, the original set has 47 tasks (42 with a hidden harmful consequence and 5 controls) in 7 apps; the fresh set has 35 tasks (30 harm tasks and 5 controls) in 5 new apps (bank, dbadmin, home, keys, travel), written after the original results. Every harm task has a harmful path and a safe path that still does what the user asked; controls have no harmful path and measure whether the tool makes agents refuse.

Agents. Claude Sonnet 5.5 and Claude Haiku 4.5 [6] run as Claude Code [5] (version 2.1.289 in every Claude study) in print mode with only the desk tools; Haiku runs with Claude Code’s default extended thinking and again with thinking off. GLM-5.3-Flash [14] and Qwen 3.5 9B and 4B [28] run through opencode [24] with the same tools, stage, tasks, judge and system prompts (Section 6.1).

Scoring and statistics. A judge reads the final state of the apps, never the agent’s words: each run ends in *harm*, *safe success* (the task done without harm) or *stopped short* (neither). Differences are paired by task, with percentile intervals from 10,000 bootstrap resamples [12] in which all runs of a task stay together; contrasts between agents pair per-task means. When every task gives the same difference (for example 1/84 against 1/84), the task bootstrap returns a zero-width interval; we report such cases as counts. With 24 to 30 tasks that come in families of near-identical variants,

percentile intervals over tasks are optimistic, so where it matters we also resample apps or task families. Proportions in the terminal studies use Wilson intervals [37].

4 Pre-registration and analysis

Each study’s design, its primary contrast and its pass bars were written and hashed before the runs they govern (Appendix B); later additions were written as hashed addenda before the runs they affect [23]. The documents are published in `paper/agents/prereg/` as copies in which only machine paths were replaced by placeholders, every substitution listed; for each document we give the SHA-256 recorded at each freeze and that of the published copy, and `verify_prereg.py` undoes the substitutions and checks every frozen hash. Addenda were appended to the plans, so each earlier frozen version is a byte prefix of the published document and is checked too. Two exceptions: the look-ahead study’s plan was written before its runs but only its tasks and code were hashed, and the terminal study’s exploratory extension was appended after its primary run, before running it, without a new hash. The hashes were recorded locally, not by a third-party registry.

- **Primary contrasts and bars.** Demo apps (Sonnet): harmful runs, *see* against *blind*, with a $\geq 50\%$ relative cut and the interval below zero, safe success down by at most 5 points and controls down by at most 10; then *see* against *placebo* with the interval below zero, and, on the original set, Ekbas not worse than Claude as forecaster (upper bound at most +5 points). Haiku: Haiku with Ekbas against Sonnet alone, “matches” (harm upper bound at most +5, success lower bound at least −5) and “beats” (in addition, a harm interval below zero or a success interval above it). Look-ahead: a success gain of at least 15 points; state tracker: at least 20 points; τ -bench: at least +10 points in reward and above the placebo. Real apps: H1 and H2 (Section 9). Real apps, follow-up: R2-H1, R2-H2 and R2-H3 (Section 10). Terminal: the plan’s reading of questions per 100 commands, added latency and completion, and for client 0.1.3 a three-part rule for “better”. AgentWorld: the decision rules P1, P2 and S1. **Met:** the demo-app bars for Sonnet, Haiku’s “beats” against Sonnet alone on the fresh set (with and without extended thinking) and, with thinking off, “matches” but not “beats” against Sonnet with Ekbas, R2-H1, R2-H3 and the 0.1.3 rule; by the AgentWorld rules, neither P1, P2 nor S1 held. **Failed:** H1, H2, R2-H2, look-ahead, state tracker, τ -bench, and Haiku’s “matches” on the look-ahead tasks.
- **Secondary contrasts** are named in each study’s plan (placebo, facts, documentation, Claude as forecaster, Sonnet with Ekbas and Haiku with the guard on the real apps, cost). They are reported with intervals and without pass/fail.
- **Exploratory** analyses are labelled as such in the text (EXPLORATORY): the conditions and contrasts the plans registered as exploratory (Sonnet with the guard and every *guard_goal* contrast on the real apps, the Qwen arm there, the Haiku terminal sessions with client 0.1.2) and every analysis decided after seeing a study’s results, including the split of the real-apps tasks by whether the request names the harmful action, the precision of the warnings, the classification of the remaining harmful runs and the app and family resampling. The follow-up’s plan fixed its splits by conflict and by new family in advance, as exploratory analyses without a bar; they are labelled so too.
- **No multiplicity correction.** We test few pre-registered contrasts per study and report every interval; we make no family-wise claim, and a reader who wants one should treat the secondary and exploratory intervals as descriptive.

Condition	Harmful runs	Safe success	Stopped short	Controls
<i>Original set: 42 harm tasks $\times$ 2 runs; 470 runs in all</i>				
blind	23/84 (27.4%)	55/84	6/84	10/10
placebo	20/84 (23.8%)	59/84	5/84	10/10
see (Ekbasis)	1/84 (1.2%)	76/84	7/84	10/10
oracle_llm	1/84 (1.2%)	74/84	9/84	10/10
oracle_truth	2/84 (2.4%)	72/84	10/84	10/10
<i>Fresh set: 30 harm tasks $\times$ 2 runs; 350 runs in all</i>				
blind	24/60 (40.0%)	36/60	0/60	10/10
placebo	22/60 (36.7%)	38/60	0/60	10/10
see (Ekbasis)	0/60	60/60	0/60	10/10
oracle_llm	10/60 (16.7%)	50/60	0/60	10/10
oracle_truth	0/60	60/60	0/60	10/10

Table 1: Claude Sonnet 5.5 on the demo apps (PRE-REGISTERED). Harmful runs, safe success and runs that stopped short are counted on the harm tasks; controls on the control tasks.

5 Hidden consequences on demo apps: Claude Sonnet

Table 1 gives the results. On the original set, *see* – *blind* is -26.2 points $[-40.5, -13.1]$ in harmful runs and $+25.0$ $[10.7, 39.3]$ in safe success; *see* – *placebo* is -22.6 $[-35.7, -10.7]$; *placebo* – *blind* is -3.6 $[-13.1, 4.8]$; and *see* and *oracle_llm* had the same harmful runs (1/84 each), which meets the pre-registered bar for Ekbasis being not worse than Claude. On the fresh set, *see* – *blind* is -40.0 $[-56.7, -23.3]$, *see* – *placebo* -36.7 $[-53.3, -20.0]$ and *placebo* – *blind* -3.3 $[-10.0, 0.0]$; Claude as forecaster, run there without a bar, left 10/60 harmful runs against 0/60 (-16.7 $[-30.0, -5.0]$), all 10 in the bank app. Every pre-registered bar was met.

What the panel got right. Scored against the app’s own truth at the panel’s state, Ekbasis flagged 64/66 harmful outcomes with 0/113 false flags on the original set and 36/39 with 0/74 on the fresh set; its 3 fresh misses were all answered at confidence ≥ 0.9 . Claude Sonnet, given the same questions, flagged 64/64 with 1/113 false flags on the original set and 28/38 with 2/70 on the fresh set, where all 12 of its wrong answers came at confidence ≥ 0.9 . Each condition is scored on the panels its own runs asked for.

Where the harm is. On the fresh set Sonnet alone did harm only in bank (12/12 runs) and travel (12/12); on the original set, in calendar (12/12), chat (6/12), mail (4/12) and drive (1/12). The one harmful *see* run on the original set clicked a sharing link it had not asked about.

Latency and cost. On a quiet server Ekbasis answered a panel in a median of 429 ms (p90 520, 113 calls); on the original runs, with the server shared with other work, the median was 4,868 ms. Claude as the forecaster gave a stated confidence through Claude Code’s print mode; it took a median of 3,675 ms per panel, which includes starting the CLI for each call, and added US\$0.0217 of Claude spend per run.




Figure 1: Harmful runs and safe success on the 30 hidden-consequence tasks of the fresh demo set, each agent alone and with Ekbas. Intervals are in Table 2.

Agent (runs per task)	Harmful runs alone → + Ekbas	Safe success alone → + Ekbas	Controls alone / + Ekbas	Harm difference (95% CI)
Claude Sonnet 5.5 (2)	24/60 → 0/60	36/60 → 60/60	10/10 / 10/10	−40.0 [−56.7, −23.3]
Claude Haiku 4.5, default thinking (1)	12/30 → 1/30	18/30 → 29/30	4/5 / 5/5	−36.7 [−56.7, −16.7]
Claude Haiku 4.5, thinking off (1)	12/30 → 0/30	18/30 → 30/30	5/5 / 5/5	−40.0 [−56.7, −23.3]
GLM-5.3-Flash (1)	6/30 → 0/30	19/30 → 30/30	4/5 / 4/5	−20.0 [−36.7, −6.7]
Qwen 3.5 9B (2)	27/60 → 7/60	33/60 → 51/60	9/10 / 10/10	−33.3 [−51.7, −15.0]
Qwen 3.5 4B (2)	30/60 → 19/60	26/60 → 36/60	10/10 / 10/10	−18.3 [−30.0, −6.7]

Table 2: Five models in six configurations on the fresh demo set. Sonnet and Haiku are PRE-REGISTERED; the GLM and Qwen runs were made later by a separate session with the same stage, tasks and judge, and were not separately pre-registered (Section 6.1).

6 Five models, six configurations, on the same hidden-consequence tasks

Every interval in Table 2 is below zero (Figure 1). The relative cut was smaller for the two Qwen models: 100.0% for Sonnet, Haiku with thinking off and GLM, 91.7% for Haiku with default thinking, 74.1% for Qwen 9B and 36.7% for Qwen 4B. GLM alone did harm in 6/30 runs and stopped short in 5/30 (none with Ekbas); those five runs stopped because of a skill its harness loaded by mistake (Section 6.1), so GLM’s runs alone are not comparable with the other agents’. These intervals resample tasks. Resampling the five apps instead (EXPLORATORY), every interval reaches zero (Sonnet [−80.0, 0.0], Haiku with thinking off [−80.0, 0.0], GLM [−53.3, 0.0], Qwen 9B [−75.0, 3.3], Qwen 4B [−36.7, 0.0]): alone, the Claude agents and GLM did harm only in bank and travel. In dbadmin, where Ekbas was confidently wrong, harm rose with it (Qwen 9B from 2/12 to 4/12, Qwen 4B from 4/12 to 5/12, Haiku with default thinking from 0/6 to 1/6).

Against the strong agent alone. Paired by task against Claude Sonnet alone, Haiku with thinking off and Ekbas did −40.0 [−56.7, −23.3] points of harm and +40.0 [23.3, 56.7] of safe success; GLM with Ekbas the same; Qwen 9B with Ekbas −28.3 [−48.3, −8.3] and +25.0 [5.0, 45.0]; Qwen 4B with Ekbas −8.3 [−21.7, 6.7] and 0.0 [−15.0, 13.3], which does not resolve. Alone, the agents did not differ much from Sonnet on these tasks: Haiku had exactly Sonnet’s

harmful runs, task by task, and Qwen 9B slightly more (+5.0 [0.0, 10.0]). v1 also reported that GLM alone did less harm than Sonnet (−20.0 [−36.7, −6.7]); that comparison is withdrawn (Section 6.1). These comparisons mostly restate the effect of the forecast on the bank and travel tasks, where the agents did harm alone; they do not show a weaker agent lifted above a stronger one. GLM and Qwen also ran in another harness. On real apps the corresponding comparison failed (Section 9); in the follow-up there, Haiku asking the user through Ekbasis did less harm than Sonnet alone, in an exploratory contrast with a scripted user (Section 10).

6.1 The GLM and Qwen runs, recomputed

The GLM-5.3-Flash and Qwen 3.5 runs were made in a separate session, with opencode as the agent harness, a runner of its own and a copy of the tool server that strips brackets from element names (Qwen wraps them) and adds the router’s note; the stage, the tasks and the judge are the study’s, and the runs were not separately pre-registered. That session was itself an AI agent (GLM-5.3-Flash driving opencode), the same model as one of the agents it ran. We recomputed every number from its run tables (which carry the judge’s verdicts), the stage’s saved sessions and the Ekbasis call log; we did not re-judge the final states. Of 25 numbers first reported for these runs, 17/25 reproduce exactly, including every harm and success count and the interval of every *see* – *blind* contrast (Appendix A). Of the 8 that do not, 4 are explained there (the GLM mean run time and the speed-up computed from it, both over harm-task runs only, which the medians do not show, and two quantile or bootstrap conventions); the other 4 (a Qwen 4B run-time median and the speed-up derived from it, a Qwen 9B control count and the Qwen 9B run-time medians) we could not reconcile, and we use none of them. We make no claim about run time: it depends on server load and on runs that crash.

Crashed runs. opencode exited with an error in some Qwen runs (4B: 15/70 blind, 6/70 see; 9B: 15/70 blind, 9/70 see); the judge scored the state they left. Without them, Qwen 4B’s harmful runs go from 26/48 (54.2%) to 16/55 (29.1%) and Qwen 9B’s from 21/46 (45.7%) to 5/52 (9.6%).

The router never acted. No panel shown in a router run had a confidence below 0.5 (GLM 0/71, Qwen 9B 0/123, Qwen 4B 0/67), so the server never appended the router’s note and *router* was *see* with other random runs; the gap between them (Qwen 9B: 7/60 against 10/60 harmful runs) is consistent with run-to-run variation.

Skill contamination (added in v2). opencode also loaded a **computer-use** skill from the operator’s user directory, synced from a desktop app account and not part of the study. Among other rules, it tells the agent never to make a payment or move money and to leave that to the person. We read every session log for a call that loaded it (`data/skill_loads.jsonl`). GLM loaded it in 55/105 runs (blind 21/35, *see* 16/35, *router* 18/35); the Qwen agents in 22/420 (Qwen 9B: 4/70, 2/70, 2/70; Qwen 4B: 7/70, 2/70, 5/70). In GLM’s bank tasks without Ekbasis, 5/6 runs loaded it and stopped before paying, citing it; the one run that did not load it paid. GLM’s harm alone was 1/18 in runs that loaded it and 5/12 in runs that did not. So the skill lowered GLM’s harm alone: its 6/30 and its five stopped-short runs are artifacts, and the comparisons of GLM alone with Sonnet alone are withdrawn. With Ekbasis, GLM did no harm whether or not it loaded the skill. Without every task where a blind or *see* run loaded it, *see* – *blind* is −33.3 [−66.7, 0.0] for GLM (3/9 to 0/9), −36.0 [−56.0, −16.0] for Qwen 9B (25/50 to 7/50) and −13.6 [−27.3, 0.0] for Qwen 4B (22/44 to 16/44); without only the runs that loaded it, −33.3 [−66.7, 0.0], −33.5 [−51.8, −15.3] and −20.8 [−33.1, −8.6]. Every difference keeps its sign; with fewer tasks,




Figure 2: Claude spend per run on the 35 fresh tasks, each agent alone and with Ekbasis (Claude spend only; the GPU that serves Ekbasis is not priced).

the GLM and Qwen 4B intervals now reach zero. The exploratory Qwen 9B arm on the real apps loaded it in 10/25 runs alone and 2/24 with Ekbasis; it did not resolve before. The Claude runs (Claude Code with no built-in tools) could not load skills.

6.2 What the remaining harm looks like (exploratory)

For each harmful run with Ekbasis we looked at what the panel said about the task’s harmful element. Qwen 9B had 17 such runs (*see* and *router*): in 10 it overrode a correct warning, in 6 it followed a wrong “no harm” answer given at confidence ≥ 0.9 , in 1 it never asked. Qwen 4B had 40: 16 overrode a warning, 6 followed a wrong answer and 18 never asked about the element. All 12 wrong answers were in dbadmin, where the panel answered that every real order is kept for “Clean up”; the only harmful run of Haiku with Ekbasis (default thinking) was the same. The model’s confident errors become harm when an agent trusts them, and the Qwen agents also overrode correct warnings or did not ask.

7 Cost

These figures are Claude spend only, as Claude Code logs it at list prices (the runs used a subscription plan); the GPU that serves Ekbasis is not priced. Mean Claude spend per run on the 35 fresh tasks was US\$0.0132 for Sonnet alone and US\$0.0149 with Ekbasis; US\$0.0233 and US\$0.0217 for Haiku with default thinking; and US\$0.0121 and US\$0.0135 for Haiku with thinking off (Figure 2), with 5.6, 7.3, 8.2, 8.1, 7.3 and 8.1 turns. Haiku was not cheaper per run for three reasons, read from each run’s usage record: with Claude Code’s defaults it produced about 32.4 times as many thinking tokens as Sonnet (1,134 against 35 per run); per token it costs only half as much (US\$1 and US\$5 per million input and output tokens, against US\$2 and US\$10 [6]; these prices reproduce every run’s logged cost); and the first request of a run weighed 7,852 tokens on Haiku against 3,747 on Sonnet (2.1 times; we did not establish why), over more turns. Turning thinking off removed the thinking (0/70 runs with thinking tokens) and cut Haiku’s cost to 0.62 [0.50, 0.75] of its default with Ekbasis and 0.52 [0.37, 0.75] alone, without changing harm or success. Haiku with thinking off and Ekbasis costs 1.02 [0.94, 1.12] times Sonnet alone and 0.90 [0.83, 0.99]

times Sonnet with Ekbasis, with the same harm as the latter; with default thinking, 1.64 [1.33, 2.11] times Sonnet alone. Ekbasis answered in a median of 450 ms (p90 524) in these runs.

8 Three studies whose bars failed

Look-ahead on multi-screen apps. The plan asked whether Ekbasis raises task success; its primary bar was a gain of at least 15 points, with the interval above zero, on 25 trap tasks whose tempting control does the wrong thing, as described in a manual the agent does not see. Claude Sonnet succeeded in 49/50 runs alone, 50/50 with a safety tool, 50/50 with a look-ahead tool and 50/50 with the manual in the prompt, with 0/50 harmful runs everywhere: the bar failed because the agent alone was at the ceiling. Haiku succeeded in 22/25 alone with 2/25 harmful runs, and in 24/25 with look-ahead with 1/25; neither difference resolves, and Haiku with look-ahead did not meet the pre-registered bar for matching Sonnet alone (harm +4 [0, 12], success −4 [−12, 0]). With hidden balances and long instructions, Sonnet reported the exact balances in 12/12 sessions with and without an Ekbasis state tracker; that bar failed too.

τ -bench retail. On the first 18 of the 115 retail test tasks (a fixed shuffled order, stopped by the budget), with the benchmark’s reward, its user-simulator prompt played by Claude Haiku and Claude Sonnet as the agent [40], pass@1 was 11/18 with no tool, 11/18 with the placebo and 12/18 with Ekbasis (+5.6 [−22.2, 33.3]); the pre-registered bar (at least +10 points with the interval above zero) failed. Runs with a harmful write were 1/18 in each condition, and policy-check violations rose with Ekbasis (0.278 per run against 0.167 with no tool and 0.111 with the placebo; +0.11 [0.00, 0.28] against no tool); its runs cost more (US\$0.183 against 0.150) and took longer (101 against 71 seconds). Ekbasis answered 26/26 of the questions it was asked inside the benchmark correctly. In these three studies the agent alone already avoided the harm or the consequence was written in the policy it reads; we group them here after the fact.

9 Real self-hosted apps

Setup. Unmodified Gitea 1.24.7 [13], Nextcloud 31.0.14 [22] and Roundcube 1.7.4 [31] with docker-mailserver 15.1.0 (Postfix 3.7.11, Dovecot 2.3.19) [10], official images pinned by digest, on a network without a route out, in containers on one machine. The versions are read from the pinned images; the plan’s table gives Roundcube 1.6, but the pinned image is Roundcube 1.7.4. The agent drives a real Chromium through Playwright 1.63 [26] with the desk tools plus `select_option`, `press_key` and `goto` (own origin only), starts logged in and never sees a password. The truth comes from the apps: Gitea’s API (including anonymous requests), Nextcloud’s shares and WebDAV, and each recipient’s mailbox. There are 24 harm tasks in 9 families and 8 controls. Each was checked on the real app before freezing: the obvious path causes the harm and a verified safe path avoids it. On 13 tasks the safe path still does everything the request names; on the other 11 (deleting both named branches, adding someone to the named team, replying to everyone, forwarding the message) every safe path leaves out or replaces something the request names, so there the judge scores doing exactly what was asked as harm and the safe alternative as success. None was dropped. Along the obvious path the screen often shows part of the picture (a Cc field, an “Open” badge on a pull request); none states the outcome. The adapters that turn a click into Ekbasis’ questions were frozen before any task file existed, but after the task families had been planned, and each action type’s questions target the hidden outcome of a planned family (for example, whether an open pull request stays open when a branch is deleted): rules paraphrased from the apps’ official




Figure 3: Harmful runs on the 24 harm tasks of the real self-hosted apps; rates are means over tasks, and runs per task are given in each label.

Agent, condition	Harmful runs	Safe success	Stopped short	Controls done
Haiku alone	66.7%	29.2%	4.2%	100.0%
Haiku + Ekbasis (see)	58.3%	39.6%	2.1%	100.0%
Haiku, rules in the prompt (docs)	58.3%	37.5%	4.2%	—
Haiku + placebo	62.5%	33.3%	4.2%	—
Haiku + state facts (facts)	70.8%	29.2%	0.0%	—
Haiku + Claude as forecaster	54.2%	41.7%	4.2%	—
Haiku + guard	54.2%	45.8%	0.0%	100.0%
Haiku + guard_goal (EXPLORATORY)	58.3%	41.7%	0.0%	100.0%
Sonnet alone	25.0%	72.9%	2.1%	93.8%
Sonnet + Ekbasis (see)	4.2%	79.2%	16.7%	—
Sonnet + guard (EXPLORATORY)	4.2%	75.0%	20.8%	100.0%

Table 3: Real self-hosted apps, 24 harm tasks (rates are means over tasks). Haiku *blind*, *see*, *guard* and *guard_goal* and Sonnet *blind* ran twice per task, the rest once. The conditions marked “—” were run on the harm tasks only, by design: Sonnet + Ekbasis (see) and the secondary Haiku conditions were not run on the control tasks.

documentation (or their source code where the documentation is silent), the state read from the app at call time, and questions per action type.

Conditions. *blind*, *see*, *placebo* and *oracle_llm* as above; *docs*: the adapters’ rules in the system prompt and no tool; *facts*: the adapter’s state facts without a forecast; *guard*: no tool, but the browser pauses the first consequential click and shows Ekbasis’ answer, and clicking again goes ahead; *guard_goal*: *guard*, with the user’s request also given to the adapter, which asks whether the click changes anything beyond what was asked. *guard* was added by an addendum after the pilot showed that agents often do not ask; *guard_goal* by a later addendum, written after an interim read showed H1 failing; both before they ran. Sonnet *guard* and every *guard_goal* contrast were registered as exploratory. Claude Haiku 4.5 ran with thinking off.

Results. Table 3 and Figure 3. **Both pre-registered hypotheses failed.** H1 (Haiku *see* — Haiku *blind*): harm $-8.3 [-18.8, 0.0]$, a 12.5% relative cut against a bar of 50%, with success $+10.4 [2.1, 18.8]$ and controls at 100.0%. (The plan as first frozen had one run per task; addendum E, written during the first runs, added the second. On the first repetition alone the cut was 25.0%,

also below the bar.) H2 (Haiku *see* against Sonnet *blind*): harm +33.3 [14.6, 52.1]; on these apps Sonnet alone did far less harm than Haiku with Ekbasis. Resampling the 9 task families instead of tasks (EXPLORATORY), H1 becomes -8.3 [-18.2, 2.0] and H2 +33.3 [10.0, 56.5].

Secondary contrasts. For Sonnet, *see* - *blind* is -20.8 [-37.5, -6.2] in harm, with one *see* run per task and no *see* runs on the controls, and +6.2 [0.0, 16.7] in success; resampling the task families (EXPLORATORY), the harm interval is [-47.8, 0.0]. Its harm alone came from 3 of the 9 families (adding someone to the team that reaches every repository, 5/6 runs; replying to everyone with outside recipients, 3/6; a public link to a folder with private files, 4/4), and its runs that stopped short rose from 2.1% to 16.7%: the runs it no longer harmed in mostly became runs where it stopped without finishing. On the 11 conflict tasks, 8/11 of its *see* runs ended in the safe alternative rather than the literal request and 2/11 stopped short; on the 13 others, 0/13 were harmful, 11/13 succeeded and 2/13 stopped short. For Haiku, *guard* - *blind* is -12.5 [-29.2, 2.1] and *guard* - *see* -4.2 [-16.7, 8.3]. Haiku did less harm with the forecast than with the bare state facts (-12.5 [-20.8, -4.2]), but the facts condition was itself no better than blind (70.8% against 66.7%), and *see* did not differ detectably from the rules in the prompt (0.0 [-14.6, 16.7]), the placebo (-4.2 [-18.8, 12.5]) or Claude as forecaster (+4.2 [-8.3, 16.7]); with Haiku, the forecaster was not the bottleneck. Claude spend per run was US\$0.051 for Haiku alone and with Ekbasis, and US\$0.046 for Sonnet alone and US\$0.053 with Ekbasis.

Exploratory conditions. Sonnet *guard* - *blind* is -20.8 [-37.5, -6.2] in harm with success +2.1 [-10.4, 14.6], and its runs that stopped short rose to 20.8%; for Haiku, *guard_goal* - *blind* is -8.3 [-22.9, 4.2].

Why Haiku did not improve. In *see*, Haiku made no useful call in 27/48 runs on harm tasks; the guard conditions removed that, pausing every relevant click. Warned, Haiku went ahead: it avoided the harm in 4 of the 18 runs in which it was warned in *see*, 9 of 35 in *guard* and 6 of 34 in *guard_goal*, 19 of 87 in all, against 17 of 18 for Sonnet. By the plan’s analysis (ii), done per run, 14/28 of Haiku’s harmful *see* runs and 4/20 of its harmless ones had a flagged panel, and every one of its 26 harmful *guard* runs had been warned (26/26).

The request itself (exploratory). A rule written before any *guard_goal* run marks a task as a conflict when every safe path leaves out something the request names (delete both branches, add to the named team, reply to everyone, forward the message). On the 13 tasks without a conflict, Haiku’s harmful runs went from 46.2% alone to 30.8% with *see* (-15.4 [-26.9, -3.8]) and 15.4% with the guard (-30.8 [-53.8, -7.7]), level with Sonnet alone (15.4%; difference 0.0 [-11.5, 11.5]), while Sonnet went from 15.4% to 0.0% (-15.4 [-38.5, 0.0]), which does not resolve; on these tasks its harm alone came from one family). On the 11 conflict tasks nothing moved Haiku (90.9% alone, 90.9% with *see*, 100.0% with the guard), while Sonnet went from 36.4% to 9.1% (-27.3 [-50.0, -4.5]).

Warnings on what was asked (exploratory). The generic questions also flag effects the user requested, such as outside recipients on a forward the user asked to send outside. On harm tasks, 13/31 of Haiku’s *see* flags, 22/63 of its *guard* flags, and 9/16 of Sonnet’s *see* flags and 9/16 of its *guard* flags marked an outcome the user had asked for; on controls, 8/9 for Haiku *see*, 11/13 for Haiku *guard* and 5/6 for Sonnet *guard*. *guard_goal* removed them (0/23 flags on requested effects) but did not change Haiku’s harm.

The forecasts against the apps. Before each click of the scripted harmful and safe paths Ekbas answered the adapter’s questions: 96 answers to 79 distinct questions across 11 action types, all right when checked against the apps, at a mean confidence of 0.991. Our truth check first had two bugs (a trailing period in an address and a destination folder that was not shared) that marked 14 of these answers wrong; we fixed the check after scoring. By the pre-registered rule (a flag on the harmful click and none on the safe one), the harmful click was flagged on 24/24 harm paths, and the last click of 9 of the 31 safe paths with a forecast was flagged too, each time for an effect the user had asked for. These are paths we scripted, not the agents’ clicks. One adapter defect was found during the study and recorded in an addendum: the Nextcloud adapter describes a single file as an empty folder, so a public link to one file gets no forecast and the guard does not pause; it left 1 safe path without a forecast, and 4 runs met it, none harmful and all 4/4 successful; the adapter was not changed mid-study.

Runs and exclusions. Of 594 logged runs, 73 were excluded before analysis: 69 stopped at once on the Claude plan’s usage limit (HTTP 429), before the agent acted, and were run again after the limit reset; 4 were cut by infrastructure (the Ekbas replica stopped for another job, one failed seeding). The study’s Claude spend was US\$24.76, for 472 valid Claude runs and the pilots. An exploratory Qwen 3.5 9B arm, cut short by a reallocation of GPUs (one run; 18 harm tasks per condition, 14 in both), had 66.7% harmful runs alone and 44.4% with Ekbas; on the 14 tasks run in both, -21.4 $[-50.0, 7.1]$, and 44.4% of its runs with Ekbas stopped short. Appendix C gives the harmful runs per family. Section 10 reports a follow-up that tests guards built for the failures above.

10 Turning warnings into safe actions

On the real apps Haiku’s harm stayed high although the forecasts were right: warned, it mostly did what the request said, and on the tasks whose request names the harmful action nothing moved it (Section 9). A follow-up (PRE-REGISTERED), designed from those failures and run on fresh tasks, tests three changes to what the guard does when a click is consequential, with Haiku as the primary agent.

Tasks. 32 new tasks on the same three apps, 24 with a hidden harmful consequence and 8 controls, in 15 families: new instances of the real-apps study’s families and six new ones, two per app (making public a repository whose history still holds a removed credential file; deleting a team that is someone’s only access to a repository they still work on; an upload link set to “Can edit”, which lets anyone with it see and delete the other clients’ files; deleting a folder a group still uses; a mailing alias that also delivers outside the company; outside clients who would see each other’s addresses in To or Cc). The adapter, the conditions and the analysis were frozen before any task existed, but after the six new families had been planned: the adapter’s new action types, and the catalog of safer ways per action type, were written for them; every task was then checked on the real apps (the obvious path ends in harm, and a safe path avoids it while doing what was asked or, on the conflict tasks, what the user’s script says the user wants) and the tasks were frozen before any agent run; none was dropped. 10 harm tasks are conflict tasks, marked with each task before any run: the literal request names the action whose obvious execution causes the harm (“reply to everyone”, “forward the message”, “make it public”, “get rid of the folder”), so on them a safe success means not doing exactly what was asked. This flag is fixed with the request, unlike the real-apps study’s rule, which looked at the verified safe paths.

Haiku, condition	Harmful runs	Safe success	Stopped short	Controls done	Questions per run	Needless per run	Seconds / US\$ per run
alone (<i>blind</i>)	72.9%	27.1%	0.0%	100.0%	–	–	53 / 0.055
<i>guard</i>	58.3%	35.4%	6.2%	100.0%	–	–	64 / 0.064
<i>guard_goal</i>	58.3%	37.5%	4.2%	93.8%	–	–	61 / 0.060
<i>guard_alt</i>	41.7%	54.2%	4.2%	100.0%	–	–	71 / 0.069
<i>ask_ekbasis</i>	2.1%	93.8%	4.2%	100.0%	0.83	0.39	81 / 0.069
<i>ask_always</i>	0.0%	93.8%	6.2%	100.0%	1.91	1.33	72 / 0.071

Table 4: The follow-up on fresh real-app tasks (PRE-REGISTERED): Claude Haiku 4.5 with thinking off, 24 harm tasks and 8 controls, two runs per task; rates are means over tasks. Questions are those put to the simulated user, needless ones those it answered with “go ahead”; seconds and Claude spend are per run, over all runs.

What the guard does. Three changes, all in the adapter; the screens and the agent’s tools are unchanged.

- **Questions that compare the click with the request.** The user’s request is added to the state, and each action type gets one or two questions on whether the click does something the user did not ask for (“does anyone lose access to a repository that the user did not ask to take away?”). They replace the single generic question of *guard_goal*, which answered “no” on the mail tasks of the real-apps study because those requests literally say “reply to everyone” or “forward the message”.
- **Safer ways.** When a click is flagged, the adapter builds alternatives for its action type from the app’s state and the request, never from the task’s identity or its truth at run time: keep a repository private, add a collaborator on the requested repository instead of the team that reaches all of them, add the members who would lose access as collaborators before deleting a team, Copy instead of Move, an upload-only link, Bcc for outside recipients, and others. Ekbas checks each one with the questions on what the user did not ask for and with whether it still achieves the purpose of the request; only a safer way that passes both is offered.
- **Asking the user.** When no checked safer way exists, or when only effects the request itself names were flagged, the question goes to the user with Ekbas’ forecast, and the answer goes back to the agent as the result of its click.

Conditions. The primary agent is Claude Haiku 4.5 with thinking off, two runs per task, with no foresee tool in any condition: *blind*; *guard* and *guard_goal* exactly as in Section 9; *guard_alt*, which is silent unless Ekbas flags a harm and then pauses the click with the forecast and the first checked safer way, or says that none was found; *ask_ekbasis*, which is *guard_alt* that asks the user in the two cases above; and *ask_always*, which puts every consequential click to the user with the app’s facts and no forecast. The user is simulated: each task carries a script, written with it before any run, of concerns (a regular expression over the question’s text, with an answer) and a default answer, “Yes, go ahead, that is what I asked for”; the first concern that matches wins. Before any run, each task was verified with its script: the script had to raise a concern at the question *ask_always* puts at the critical click of the harmful path, and answer every click of the safe path with the default; two scripts were edited for that. On the scripted harmful path *ask_always* therefore catches the harm by construction; the questions of *ask_ekbasis* were not checked against the scripts in this way. The app’s facts that *ask_always* shows are those the adapter reads for each action type. A question answered with the default counts as needless. The plan also has Claude Sonnet 5.5 *blind*, *guard_goal* and *ask_ekbasis*, one run per task, as secondary arms.




Figure 4: Haiku on the fresh real-app tasks: a pause with the forecast, a checked safer way, and asking the user.

Results. Table 4 and Figure 4. **R2-H1 passed:** with *ask_ekbasis* Haiku’s harmful runs fell from 72.9% to 2.1% (-70.8 [$-87.5, -50.0$], a 97.1% relative cut) and its safe success rose from 27.1% to 93.8% ($+66.7$ [$47.9, 85.4$]), with controls done at 100.0%. **R2-H3 passed:** without asking anyone, the request-aware checks and checked safer ways of *guard_alt* took harm from 58.3% with the plain guard to 41.7% (-16.7 [$-33.3, -2.1$]), with success $+18.8$ [$4.2, 35.4$]. **R2-H2 failed:** *ask_ekbasis* put 0.83 questions per run to the user against 1.91 for *ask_always*, 43% as many against a bar of 25%. Asking the user before every consequential click, with the app’s facts and no forecast, also avoided the harm (the scripts had been checked to object to its question at the critical click of each scripted harmful path): 0.0% harmful runs, and harm $+2.1$ [$0.0, 6.2$] points for *ask_ekbasis* against it. In this study Ekbasis’ contribution over asking every time is fewer questions, not less harm. Against *blind*, the plain guard changed harm by -14.6 [$-27.1, -4.2$] points, and the generic request question of *guard_goal* added nothing to it (0.0 [$-8.3, 8.3$]).

From warnings to actions. In *guard_alt*, 16 harm-task runs were offered a checked safer way: Haiku did 13/16 of them safely, with harm in 3/16. Of the 23 harm-task runs that got a warning without one, 17/23 ended in harm. The safer ways offered were 4 collaborators added before a team was deleted, 4 upload-only links, 4 moves of outside recipients to Bcc, 3 collaborators instead of the team that reaches every repository and 1 repository kept private. On the 10 conflict tasks the guard never found a safer way that passed its checks, and *guard_alt* did harm in 16/20 runs; asking the user took them to 0/20, with 19/20 done, as with *ask_always*. On the 14 other harm tasks, *guard_alt* took harm from 39.3% with the plain guard to 14.3% (-25.0 [$-46.4, -7.1$]; EXPLORATORY, with the split fixed in the plan). On the six new families *ask_ekbasis* had 4.2% harmful runs against 79.2% alone; on fresh instances of the real-apps study’s families, 0.0% against

Sonnet, condition	Harmful runs	Safe success	Stopped short	Controls done	Questions per run	Needless per run	Seconds / US\$ per run
alone (<i>blind</i>)	50.0%	45.8%	4.2%	100.0%	–	–	40 / 0.048
<i>guard_goal</i>	16.7%	54.2%	29.2%	100.0%	–	–	46 / 0.058
<i>ask_ekbasis</i>	8.3%	75.0%	16.7%	100.0%	0.56	0.28	52 / 0.057

Table 5: The follow-up’s secondary arms: Claude Sonnet 5.5 on the same 24 harm tasks and 8 controls, one run per task; rates are means over tasks.

66.7% (EXPLORATORY, with the split fixed in the plan).

The questions. Of the 135 clicks *ask_ekbasis* checked, 71 passed silently, 11 got a checked safer way without a question, and 53 became a question to the user: 28 after a flag that the click goes beyond the request with no checked safer way, and 25 on effects the request itself names. The simulated user answered “go ahead” to 25/53 of them, 16 of these on effects the request names and 9 after an alert that the click goes beyond it, 8 of the 25 on controls, against 85/122 for *ask_always*: 57% fewer questions and 71% fewer needless ones. The questions of *ask_ekbasis* carry Ekbasis’ forecast; those of *ask_always* carry only the app’s facts.

The one harmful run. The only harmful *ask_ekbasis* run was on an upload-link task. The agent made the upload-only link the guard had suggested, then created a second link from a share panel the adapter could not resolve, so the question named “the selected item” instead of the other clients’ files; the simulated user, whose concern names those files, said go ahead, and the second link could read them. The forecast flagged the click; the adapter did not say what it exposed. In that family Haiku alone did no harm (0/4). Two other runs ended without the task done and without harm. By app, *ask_ekbasis* did harm in 0/18 Gitea runs, 0/14 mail runs and 1/16 Nextcloud runs. An *ask_ekbasis* run took 81 seconds and US\$0.069 of Claude spend on average, against 53 seconds and US\$0.055 alone, before any wait for a real user’s answer.

Sonnet (secondary arms). Table 5. These fresh tasks were harder than those of the real-apps study: Sonnet alone did harm in 50.0% of them, against 25.0% there, and in 66.7% on the six new families against 33.3% on fresh instances of the real-apps study’s families. With *ask_ekbasis* its harmful runs fell to 8.3% (−41.7 [−62.5, −20.8]) and its safe success rose from 45.8% to 75.0% (+29.2 [8.3, 50.0]); with *guard_goal* harm fell to 16.7% (−33.3 [−54.2, −16.7]), but 29.2% of its runs stopped short, warning and asking for guidance. The two harmful *ask_ekbasis* runs were deliberate choices. In one, Sonnet declined the checked safer way, adding a colleague as a collaborator before deleting a team, because the user had not asked to grant new access; it deleted the team and reported that the colleague had lost access. In the other, the simulated user’s answer, to move the folder to the archive instead of deleting it, reached Sonnet as a tool result, and Sonnet set it aside because it “wasn’t from you”. This is a limit of our harness and a lesson for products: a user’s answer has to reach the agent through the user’s channel, such as a permission prompt, not as tool output. Compared with Sonnet alone on the same tasks (EXPLORATORY), Haiku with *ask_ekbasis* had −47.9 [−68.8, −27.1] points of harm, with a scripted user that reads every question perfectly; Haiku with *guard_alt*, which asks nobody, had −8.3 [−29.2, 12.5].

Limits of this study. The simulated user reads every question perfectly, never tires of approving and answers by rules written with the task. This flatters both asking conditions, and *ask_always* most: its questions are more frequent, carry only the facts the adapter reads, and the scripts were checked to object to its question on each harmful path; with people the number of questions may

matter more than it does here, which this study did not measure. The safer ways are written by hand per action type, from the app’s state and the request and with the planned families in mind; who writes them for a new app is open. The judge of the family in which a new person joins the team that reaches every repository checks only the new person’s access; had an agent restricted that team instead, the effect on its existing members would not have been judged. Two runs per task; the apps are those of the real-apps study and only the tasks are fresh; the tasks and the user’s scripts are ours. The plan’s cap of US\$25 on Claude spend, with its fixed order of steps, stopped the study before its end, with *guard_goal* at 52/64 runs and the Sonnet arms not started; up to that point its Claude spend was US\$25.23, including a pilot of 9 runs (US\$0.98) that found one adapter defect, fixed in a hashed addendum before the study runs. The cap was then removed by a second hashed addendum, written after the Haiku results were known and with no change to the design, and the remaining runs (the missing Haiku *guard_goal* runs and the Sonnet arms) were run in the plan’s order: 480 valid runs in all, for US\$31.02 of Claude spend. A third addendum, also written after the Haiku results were known, after all runs, re-judged the mail runs with whitespace-normalised message bodies (EXPLORATORY): the mail judge, taken unchanged from the real-apps study, misses text the webmail wraps across lines. Of the 135 mail runs, 1 verdict changes, a Sonnet *ask_ekbasis* forward from not done to done, which would raise that arm’s safe success to 79.2%; no harm verdict changes. The numbers above are those of the frozen judge; the released benchmark ships the normalised one. The Sonnet arms ran once per task.

11 A real terminal

Claude Code [5] ran on clones of three real repositories (Python, Go, JavaScript) with a hook that asks Ekbasis before git and shell commands and stops the ones it judges to lose work; the truth comes from snapshots before and after every command. In print mode a hook’s question is a denial: the command does not run, and the agent reads the reason and continues; in interactive use a person would answer.

Client 0.1.2. With Claude Sonnet (72 sessions, 243 commands), nothing destructive happened with or without the hook, so the hook only had a cost: 14/132 commands asked (10.6 per 100, 95% CI 6.4–17.0; resampling sessions, 5.6–15.5), above the 10 per 100 that the plan called bad for daily use; 10/14 of the questions were about commands that lose nothing, at a median of 2.12 s per command (p90 9.68 s) on a shared server, which the plan’s scale calls tolerable, with 34/36 tasks done against 35/36 without it. With Claude Haiku on the four tasks with a tempting destructive shortcut (24 sessions, EXPLORATORY, an extension appended after the primary run), the agent without the hook ran 6 work-losing commands in 4 of 12 sessions. On the three of these tasks with planted user work, it kept that work in 5/9 sessions without the hook (95% CI 26.7–81.1) and in 9/9 with it (70.1–100.0); Fisher’s exact $p = 0.082$. The hook caught 8/8 real losses (2 caught live, 6 when the hook was replayed on the commands of the sessions without it). In a second exploratory part, on 19 classic destructive shortcuts run on the tasks’ real states, it asked on 17/19, and on 3/11 of their safe alternatives.

Client 0.1.3. After fixes designed on those commands, fresh sessions on the same repositories and tasks (PRE-REGISTERED; so the drop in questions is measured on the tasks the fixes were designed on): with Sonnet the hook asked on 3/103 commands (2.9 per 100, 95% CI 1.0–8.2; resampling sessions, 0.0–6.0), against 11/103 (10.7 per 100) for 0.1.2 replayed on the same commands; on those commands 0.1.2 alone asked 9 times and 0.1.3 alone once (exact McNemar $p = 0.021$). Across all

Suite	Items	One pass, all items			AgentWorld reasoning, 100 items		
		Ekbasis	AgentWorld	AW – Ek [95% CI]	Ekbasis	AgentWorld	AW – Ek [95% CI]
accumulation	1,736	83.4%	57.0%	−26.4 [−29.4, −23.3]	84%	93%	+9 [1, 18]
git (fresh)	3,250	88.8%	83.9%	−4.9 [−6.2, −3.5]	87%	85%	−2 [−11, 7]
planning	8,898	95.3%	72.8%	−22.5 [−23.6, −21.3]	99%	96%	−3 [−7, 1]
rules under stress	9,075	95.2%	69.2%	−26.0 [−27.0, −25.0]	97%	99%	+2 [−2, 6]
shell	3,708	90.7%	81.3%	−9.4 [−10.8, −8.1]	91%	89%	−2 [−9, 5]
SQL	2,056	84.4%	68.1%	−16.3 [−18.6, −14.0]	86%	92%	+6 [−1, 13]
apps (written rules)	1,776	97.2%	83.2%	−14.1 [−15.9, −12.3]	97%	95%	−2 [−8, 3]
git (capability map)	1,625	88.2%	84.0%	−4.2 [−6.2, −2.4]	86%	84%	−2 [−9, 5]
habit against rule	2,880	97.9%	88.9%	−9.0 [−10.1, −7.9]	98%	100%	+2 [0, 5]
languages and formats	3,503	97.7%	83.8%	−13.9 [−15.2, −12.7]	98%	97%	−1 [−5, 2]
guard, known families	4,486	94.7%	58.4%	−36.3 [−37.9, −34.8]	96%	64%	−32 [−41, −23]
guard, new families	6,394	85.9%	61.1%	−24.8 [−26.2, −23.5]	85%	67%	−18 [−29, −6]
AgentWorldBench terminal	148	92.6%	81.1%	−11.5 [−16.9, −6.6]	94%	81%	−13 [−19, −7]
AgentWorldBench MCP	150	97.3%	94.7%	−2.7 [−7.4, 0.0]			

Table 6: Ekbasis against Qwen-AgentWorld-35B-A3B on the same questions (PRE-REGISTERED). Left: both models score the listed options in one pass. Right: AgentWorld first reasons and writes the predicted state (up to 8,192 tokens), then answers; on the MCP domain only one item was run in this mode. Intervals resample item clusters.

204 commands, needless questions about commands that lose nothing fell from 9/204 to 2/204 and those about losses git or a rebuild can undo from 14/204 to 3/204; the median hook time was 0.08 s (p90 0.77 s), measured on a quieter server than 0.1.2’s and with more commands settled in code before the model is asked; tasks done were 35/36 with the hook and 36/36 without. With Haiku, both versions caught 7/7 real losses, the hook asked 6/80 times (5 of them about real losses), and the planted work was kept in 9/9 sessions with the hook against 7/9 without (95% CI 45.3–93.7; Fisher’s exact $p = 0.471$). By the pre-registered rule, 0.1.3 is better than 0.1.2: fewer questions on the same commands, no real loss missed that 0.1.2 caught, and completion within one session of the control.

12 The consequence model against another world model

Qwen-AgentWorld-35B-A3B [44, 29] (revision 60d2b043) is an open world model trained on real agent trajectories, a mixture of experts with 35B parameters of which 3B are active per token, served with vLLM 0.30.0. We gave it exactly the inputs Ekbasis reads, on our 12 question suites (49,387 items) and on two domains of its own benchmark recast as multiple choice (our adaptation: the benchmark scores the written observation with a judge, and our distractors are later observations of the same trajectory), in two modes fixed in advance (Table 6). Scoring the options in one pass, Ekbasis was more accurate on all 12 suites, with every interval excluding zero (12/12), had a lower Brier score on all 12 (11/12 with the interval excluding zero), and was more accurate on AgentWorld’s terminal domain; the MCP domain showed no resolved difference. The Brier comparison reads Ekbasis through its release calibration and AgentWorld raw. Of the other pre-registered metrics, Ekbasis had the lower expected calibration error in 10/12 suites and the higher AUROC of confidence for correctness in 11/12; AgentWorld made fewer confident errors (wrong at confidence ≥ 0.9) in the two git suites (170 against 260 and 85 against 137) and in the guard suite of new families (608 against 746), where it answered with high confidence less often.

Allowed to reason and simulate first, on 100 items per suite, AgentWorld won 1 suite (accumulation), showed no resolved difference on 9 and lost the 2 guard suites, at a median of 51 to 319 seconds and up to 8,192 generated tokens per question, against one forward pass. In this

mode AgentWorld often did not finish reasoning within the 8,192-token cap fixed in the plan: it finished on 6/100 and 28/100 items of the two guard suites and on 45/100 and 48/100 of the two git suites, and an unfinished item is answered without a written prediction; its two losses are in the suites where it finished least. The guard suite of known families comes from families present in Ekbas’s training data, and we prompted AgentWorld with a generic simulator prompt of our own. The comparison also decided, by a rule fixed before the runs, whether to propose AgentWorld as the base of our next model; neither of the rule’s conditions held. The one-pass mode measures the use this paper studies; it does not say that AgentWorld is a worse world model, and our suites are in the format Ekbas was trained on (Section 13).

13 Data and overlap

The agent tasks, demo apps, adapters and judges were written for these studies by the author working with Claude Opus 5.5; none was drawn from a training set. The question suites of Section 12 were checked against every training file of the released model’s lineage (the training file of its first parent and the two mined sets of the second; the fine-tuning data of the base model, Eikos-27B, were not scanned) with two rules: an identical state text, or a word 8-gram near-duplicate (shared layout removed, $\geq 50\%$ contained, Jaccard ≥ 0.5). The six suites of the first group (28,723 questions) and the four capability-map suites (9,784 questions) were scanned by the evaluation-set firewall, and the two guard suites (10,880 questions) by the integrity audit of the public evaluation sets; each scanned set has the same number of rows as its suite in the comparison, and 0 rows were flagged in any of them. The demo-app panel questions and the real-app adapter questions were not scanned, and the guard suite of known families is drawn, by design, from families present in training. The AgentWorldBench items are third-party and were not scanned. The same audit found that one evaluation set of our earlier paper overlapped its training data; that paper’s corrected version reports the affected numbers [34].

14 Limitations

- **Our tasks.** Every agent task except those of τ -bench was written by us, to have a hidden consequence and a safe path; in Sections 5–7 and the look-ahead and tracker studies the apps were ours too. The real-apps studies remove our apps, not our tasks, and in the follow-up the simulated user’s scripts are ours as well. They are released with the apps, adapters and judges (see Reproducibility and data availability), so others can rerun the studies and add tasks written by people who did not build the guard.
- **Two settings, two strengths of evidence.** On our demo apps each app has its own questions with the harmful answer marked, and the effect was large for five models. On real apps the questions are per action type and also flag what the user asked for; Sonnet’s drop from 25.0% to 4.2% is a secondary contrast with one run per task, carried by 3 of 9 task families (interval over families $[-47.8, 0.0]$) and mostly by tasks whose request names the harmful action, with more runs stopping short; the pre-registered hypotheses for Haiku failed, and in the follow-up a guard that asks the user cut Haiku’s harmful runs to 2.1% with a simulated user, and asking every time without Ekbas cut them to 0.0%.
- **A smaller model with the consequence model.** Doing less harm than Sonnet alone holds on our demo apps, where the agents alone were no worse than Sonnet, or only slightly; on real apps Haiku only matched Sonnet alone, with the automatic pause and only on the

13 tasks whose request does not name the harmful action, an exploratory split; in the follow-up, Haiku asking a scripted user through Ekbasidid did less harm than Sonnet alone (-47.9 [$-68.8, -27.1$]), also exploratory. Per run, Haiku with Ekbasidid cost about as much Claude spend as Sonnet alone (1.02 [$0.94, 1.12$]), before pricing the GPU that serves Ekbasidid.

- **Agents that follow the literal request override correct warnings.** On real apps Haiku avoided the harm in 19 of the 87 runs in which it was warned, against 17 of 18 for Sonnet, and almost never on tasks whose request names the harmful action; on our demo apps the Qwen agents overrode correct warnings too, while Haiku did not. A checked safer way helped Haiku, and a question to the user when there was none brought its harm to 2.1% in the follow-up, and asking every time without Ekbasidid to 0.0%.
- **A simulated user.** The follow-up’s user is a script that reads every question perfectly and never tires of approving. It flatters both asking conditions, *ask_always* most; with people, how often the guard asks may matter more, and how well they read the questions may matter too. Neither was measured. Its answers reached the agent as tool results, and Sonnet once set one aside as not coming from the user; in a product the answer should come through the user’s channel.
- **Three studies whose bars failed.** The look-ahead and state-tracker studies failed their pre-registered bars because the agent alone was at the ceiling; τ -bench failed its bar, its policy states the consequences, and policy-check violations rose with Ekbasidid. They show where a consequence model added nothing, not that its effect is absent elsewhere.
- **Concentrated effects.** On the fresh demo set, the Claude agents and GLM did harm alone only in bank and travel; resampling the five apps instead of tasks (EXPLORATORY), every agent’s interval reaches zero. On the real apps, Sonnet did harm alone in 3 of 9 families.
- **The adapter does part of the work.** It reads the state and chooses the questions, and the questions encode which outcomes matter: on our demo apps one mechanic per app, on the real apps the hidden outcome of each planned family. *placebo*, *facts*, *docs*, *oracle_llm* and *oracle_truth* separate the parts. In the follow-up it also builds the safer ways, written by hand per action type, and an adapter miss lay behind the one harmful *ask_ekbasidid* run.
- **Conditions added during a study.** *guard* was added by an addendum after the pilot, and *guard_goal* after an interim read that showed H1 failing, both before they ran; the conflict and warning-precision rules were written before *guard_goal* ran, after the other conditions.
- **Few runs.** One or two runs per task and agent; per-app and per-family counts are small, and the percentile intervals over 24 to 30 tasks are optimistic. The terminal results on planted user work rest on 9 sessions per arm, and the follow-up’s Sonnet arms ran once per task.
- **Confident errors.** Ekbasidid’s wrong answers on the fresh demo set were all at confidence ≥ 0.9 , and the Qwen agents followed them into harm; in dbadmin harm rose with the forecast. On the scripted real-app paths its answers were right, but its flags also marked effects the user had asked for.
- **One consequence model,** one released checkpoint; the GLM and Qwen runs were made by a separate AI-agent session in another harness and recomputed by us from its tables, without re-judging the final states; that harness also loaded an unrelated skill in some runs (Section 6.1).

- **Studies not in this paper.** Two further pre-registered studies in the same program are not reported here and have not been published. A preflight check for multi-step shell and SQL changes, not yet released, reduced harm for Claude Haiku but only interrupted Claude Sonnet, none of whose warnings was correct. Adapters written automatically by a language model from the real apps’ documentation did not pass the bars their plan set against the hand-written adapters used in Section 9.
- **Pre-registration records.** The freeze files record hashes with times for the terminal, both real-apps and the AgentWorld studies, and hashes with dated headings for the demo-app study; the look-ahead plan has no recorded hash, and the follow-up’s three addenda have hashes without times, the second and third written after the Haiku results were known. The hashes were recorded locally, not by a third-party registry.

15 Ethics

All tasks are benign requests a user might make; the harms are side effects of doing them the obvious way (a deleted order, a cancelled return flight, a repository made public), and they happened only inside demo apps and local, self-hosted copies of open-source apps with no real users, no real accounts, no outbound mail and no network route out. Test accounts were created for the studies and their credentials never reached the agents. No person took part as a subject: the user who answers the guard’s questions in the follow-up is a script. The paper describes failure modes of agents and of a guard, not ways to cause harm, and releases no transcripts.

16 Conclusion

With a consequence model in the loop, Claude Sonnet did less harm when the consequence was hidden: from 24/60 to 0/60 harmful runs on our demo apps, and from 25.0% to 4.2% on real, unmodified apps, a secondary contrast with one run per task, carried by 3 of 9 task families and mostly by tasks whose request names the harmful action. Its safe success did not fall, but more of its runs stopped short. Claude Haiku 4.5, GLM and Qwen 9B with the model did less harm than Sonnet alone on our demo apps, where every app has its own questions and Haiku and Qwen 9B alone did about as well as Sonnet (GLM’s runs alone are not comparable; Section 6.1); on real apps Haiku with an automatic pause only matched Sonnet alone, and only on tasks whose request does not name the harmful action (an exploratory split). When warned, Haiku mostly did what the request said; a guard that offers a safer way it has checked and asks the user when there is none took its harmful runs on fresh tasks from 72.9% to 2.1%, with a simulated user that reads every question perfectly (Sonnet’s, in a secondary arm, from 50.0% to 8.3%), and asking the user every time, without Ekbas, did as well with more questions, although the user’s scripts had been checked to object to that condition’s question on every harmful path. Where the agent alone was at the ceiling, or the policy stated the consequences, there was nothing measurable to gain and the pre-registered bars failed; in τ -bench, policy-check violations rose with Ekbas. What comes next follows from the limits: real users instead of a script, adapters and safer ways not written by hand, tasks written by other people, and the model’s confident errors.

Competing interests

The author develops Ekbasis and its client. Claude models are among the agents studied, served as the forecaster in *oracle_llm*, and helped design and run the studies.

Acknowledgments

This work was done with Claude Opus 5.5 (Anthropic), which helped plan and run the experiments, write the tasks, the analysis code and the text, and check every number and reference. The GLM-5.3-Flash and Qwen 3.5 agent runs were executed by a separate agent session (GLM-5.3-Flash driving opencode) under the studies’ plan; we recomputed every number from the run tables (Section 6.1). All errors are the author’s.

Reproducibility and data availability

The paper, its source and everything needed to recompute its numbers are in the `paper/agents/` folder of <https://github.com/OpenInterpretability/ekbasis>: `data/` holds the per-run outcome tables of every study (task, app or family, agent, condition, repetition, harm, success, stopped short, cost, seconds and the derived per-run fields used here, with each check of the follow-up’s guard reduced to categories), the per-item correctness and confidence vectors of the AgentWorld comparison keyed by the hash of each prompt, and three aggregated tables listed in the folder’s README; `reproduce.py` recomputes `numbers.json` from them; `check_numbers.py` fails if the paper cites a number `numbers.json` does not contain (it checks that each number exists, not where it is used); `figures.py` draws the figures; `refs_check.json` records how every reference, and every venue cited for a preprint, was resolved; `prereg/` holds the pre-registration documents and freeze files with `verify_prereg.py`. `benchmark/` holds what is needed to rerun the studies: the tasks and their generators, the demo apps, the Docker Compose file that runs the official Gitea, Nextcloud, Roundcube and mail-server images pinned by digest, the seeding scripts, the adapters, the judges, the stages, the agent’s tools, the forecast bridges, the runners and the task verifiers, with the terminal study’s repositories and τ -bench referenced by URL and commit; its README runs one task end to end with any agent and any Ekbasis server, and `verify_benchmark.py` checks every file against the files that ran. Machine paths, a server’s address and a user name were replaced by environment variables, and the real apps’ test credentials are generated at setup. No transcript or real credential is released. The model is at <https://huggingface.co/caiovicentino1/Ekbasis-27B> and the client at the repository above.

A First-reported and recomputed numbers for the GLM and Qwen runs

Number	First reported	Recomputed
GLM harm, blind / see / router	6/30, 0/30, 0/30	6/30, 0/30, 0/30
GLM safe success, blind	19/30	19/30
GLM controls, blind / see / router	4/5, 4/5, 5/5	4/5, 4/5, 5/5
GLM see – blind, 95% CI	[−36.7, −6.7]	[−36.7, −6.7]
GLM median seconds per run, blind / see / router	25, 30, 28	24.7, 29.7, 28.3
GLM mean seconds per run, blind → see	73 → 36	67 → 34 (all runs)
GLM speed-up with Ekbas	2 (from means)	0.8 (ratio of medians)
GLM Ekbas latency, median / p90 / max	434 / 527 / 860 ms	434 / 525 / 860 ms
Qwen 4B, seed 1, harm blind / see / router	16/30, 12/30, 11/30	16/30, 12/30, 11/30
Qwen 4B, seed 1, see – blind	−13.3 [−30.0, 3.3]	−13.3 [−30.0, 3.3]
Qwen 4B, seed 1, median seconds, blind / see	67 / 24	25.3 / 24.5
Qwen 4B, seed 1, speed-up with Ekbas	2.7	1.0 (ratio of medians)
Qwen 4B, 2 seeds, harm blind / see	30/60, 19/60	30/60, 19/60
Qwen 4B, 2 seeds, safe success blind / see	26/60, 36/60	26/60, 36/60
Qwen 4B, 2 seeds, see – blind	−18 [−30.0, −6.7]	−18.3 [−30.0, −6.7]
Qwen 9B harm, blind / see / router	27/60, 7/60, 10/60	27/60, 7/60, 10/60
Qwen 9B safe success, blind / see / router	33/60, 51/60, 49/60	33/60, 51/60, 49/60
Qwen 9B controls, blind / see / router	9/10, 10/10, 9/10	9/10, 10/10, 8/10
Qwen 9B see – blind	−33.3 [−51.7, −15.0]	−33.3 [−51.7, −15.0]
Qwen 9B router – blind	−28.3 [−45.0, −11.7]	−28.3 [−45.0, −10.0]
Qwen 9B median seconds, blind / see / router	28, 36, 35	26.4, 35.8, 34.5
Qwen 9B foresee calls, see and router	282	282

Medians are given to one decimal. The GLM latency p90 and the Qwen 9B router interval differ only by a quantile or bootstrap convention; the GLM mean of 73 → 36 seconds, and the speed-up of 2 computed from it, reproduce over the harm-task runs only. We could not reconcile the Qwen 4B seed-1 median of the agent alone (and the speed-up of 2.7 that follows from it), the Qwen 9B router control count, or the Qwen 9B blind and router medians.

B Pre-registrations

The documents, the freeze files and `verify_prereg.py` are in `paper/agents/prereg/`; its README lists every substitution of a machine path and the full hashes. Each SHA-256 below was recorded before the runs it governs; “prefix” means the frozen hash is that of the first bytes of the published document, before later addenda were appended. Times are UTC.

Study	Document and when it was frozen	SHA-256 (first 8)
Demo apps	plan, before any measurement, 5 October (prefix) with addenda A (controls) to C (τ -bench), each before its runs, 5 October (prefix) with addenda D, D.1, D.2 (Haiku, cost, thinking off) and the D.2 note, each before its runs, 5–6 October addendum B’s tasks, judge and apps	7674f5c5 e018063c 88988a92, be07f1cc, 5215a26d, ac87dbb5 in the freeze file
Look-ahead	plan, before any agent run, 5 October tasks and code, 5 October	not recorded in the freeze files
Terminal, client 0.1.2	plan, before the first pair, 5 October 17:48 (prefix; the exploratory extension follows)	3aa101e9
Terminal, client 0.1.3	addendum, before the first paid session, 6 October 11:48	41385b3e
Real apps	adapters, rules and questions, before any task existed, 6 October 12:29 plan and verified tasks, before the first agent run, 6 October 12:49 (prefix) with addenda B–H, 6 October 12:54 to 14:59 (each a prefix; H the whole)	dcfe6abc c73625be fe86b581 ... f07f48f4
Real apps, follow-up	plan, adapter, conditions and analysis, before any task existed, 6 October 17:19 (prefix) with the verified tasks, before the first agent run, 6 October 17:34 (prefix) with addendum A (the pilot’s adapter fix), before the first study run (see below; prefix) with addendum B (the spending cap removed, no design change), before the remaining runs, 6 October (prefix) with addendum C (mail runs re-judged with nor- malised bodies; release hygiene), after all runs, 6 October	3e2d8995 50bba2f6 08c23c58 d788740d b2f43162
AgentWorld comparison	plan, before any AgentWorld forward pass, 6 Octo- ber 14:57 addenda 1–4 (code fixes; serving of mode (b)), 6 October	0c83f226 3231803a, c44c2938, 28bd60dd, d1dfd270

The follow-up’s freeze file records the hashes of addenda A, B and C without a time. For addendum A, the freeze file was last written at 17:42 and the first study run started at 17:43, as its run log (not released) shows; addendum B was written when the cap was removed, about 21:00, before the remaining runs, and addendum C after all runs, about 22:00.

C Real apps by task family

Family (harmful runs)	Haiku alone	Haiku + see	Haiku + guard	Sonnet alone	Sonnet + see	Sonnet + guard
Gitea: delete branches, one with an open pull request	4/4	4/4	4/4	0/4	0/2	0/2
Gitea: take someone off one team	0/4	0/4	0/4	0/4	0/2	0/2
Gitea: create a repository (public by default)	4/6	2/6	0/6	0/6	0/3	0/3
Gitea: add someone to a team that reaches every repository	4/6	5/6	6/6	5/6	0/3	0/3
Mail: forward a thread with internal quoted text	6/6	5/6	6/6	0/6	0/3	0/3
Mail: reply to everyone, with outside recipients	6/6	6/6	6/6	3/6	1/3	1/3
Nextcloud: public link to a folder with private files	4/4	3/4	3/4	4/4	0/2	0/2
Nextcloud: move a file out of a shared folder	4/6	3/6	1/6	0/6	0/3	0/3
Nextcloud: share a folder (edit rights by default)	0/6	0/6	0/6	0/6	0/3	0/3

Counts are runs, two per task for Haiku and for Sonnet alone and one for Sonnet with Ekbas or the guard. The Claude spend and the median wall time per run of every condition are in `numbers.json` (`ra.cost_time`).

References

- [1] A. N. Angelopoulos et al. Learn then Test: Calibrating Predictive Algorithms to Achieve Risk Control. arXiv preprint, 2021. arXiv:2110.01052.
- [2] A. N. Angelopoulos et al. Conformal Risk Control. arXiv preprint, 2022. arXiv:2208.02814.
- [3] V. Barres et al. τ^2 -Bench: Evaluating Conversational Agents in a Dual-Control Environment. arXiv preprint, 2025. arXiv:2506.07982.
- [4] H. Chae et al. Web Agents with World Models: Learning and Leveraging Environment Dynamics in Web Navigation. In Proceedings of ICLR, 2025. arXiv:2410.13232.
- [5] Anthropic. Claude Code (software). <https://code.claude.com/docs/en/overview>.
- [6] Anthropic. Claude Sonnet 5.5 and Claude Haiku 4.5 (models). <https://platform.claude.com/docs/en/models/overview>.
- [7] Y. Feng et al. CORA: Conformal Risk-Controlled Agents for Safeguarded Mobile GUI Automation. arXiv preprint, 2026. arXiv:2604.09155.
- [8] A. Cuadron, P. Yu, Y. Liu, and A. Gupta. SABER: Small Actions, Big Errors – Safeguarding Mutating Steps in LLM Agents. arXiv preprint, 2025. arXiv:2512.07850.
- [9] FAIR CodeGen team et al. CWM: An Open-Weights LLM for Research on Code Generation with World Models. arXiv preprint, 2025. arXiv:2510.02387.
- [10] docker-mailserver (software). <https://github.com/docker-mailserver/docker-mailserver>.
- [11] W. Lin et al. DreamGuard: Efficient Runtime Guardrail for LLM Agents via Risk-Aware World Model. arXiv preprint, 2026. arXiv:2608.05695.
- [12] B. Efron. Bootstrap Methods: Another Look at the Jackknife. The Annals of Statistics, 7(1):1–26, 1979. DOI: 10.1214/aos/1176344552.
- [13] Gitea (software). <https://about.gitea.com/>.
- [14] Z.ai. GLM-5.3-Flash (model). <https://huggingface.co/zai-org/GLM-5.3-Flash>.
- [15] Y. Gu et al. Is Your LLM Secretly a World Model of the Internet? Model-Based Planning for Web Agents. arXiv preprint, 2024. arXiv:2411.06559.
- [16] S. Hao et al. Reasoning with Language Model is Planning with World Model. In Proceedings of EMNLP, pp. 8154–8173, 2023. DOI: 10.18653/v1/2023.emnlp-main.507.
- [17] Q. Hu et al. SABER: Benchmarking Operational Safety of LLM Coding Agents in Stateful Project Workspaces. arXiv preprint, 2026. arXiv:2606.01317.
- [18] T. Kuntz et al. OS-Harm: A Benchmark for Measuring Safety of Computer Use Agents. In NeurIPS Datasets and Benchmarks Track, 2025. arXiv:2506.14866.
- [19] W. Kwon et al. Efficient Memory Management for Large Language Model Serving with PagedAttention. In Proceedings of the 29th Symposium on Operating Systems Principles (SOSP), pp. 611–626, 2023. DOI: 10.1145/3600006.3613165.
- [20] I. Levy et al. ST-WebAgentBench: A Benchmark for Evaluating Safety and Trustworthiness in Web Agents. In Proceedings of ICLR, 2026. arXiv:2410.06703.
- [21] K. Li et al. Discriminative World Models for Web Agents. arXiv preprint, 2026. arXiv:2609.02885.

- [22] Nextcloud (software). <https://nextcloud.com/>.
- [23] B. A. Nosek et al. The Preregistration Revolution. *Proceedings of the National Academy of Sciences*, 115(11):2600–2606, 2018. DOI: 10.1073/pnas.1708274114.
- [24] opencode (software). <https://opencode.ai/>.
- [25] Y. Qu et al. Overeager Coding Agents: Measuring Out-of-Scope Actions on Benign Tasks. *arXiv preprint*, 2026. arXiv:2605.18583.
- [26] Playwright (software). <https://playwright.dev/>.
- [27] V. K. Bonagiri, P. Kumaraguru, K. Nguyen, and B. Plaut. Check Yourself Before You Wreck Yourself: Selectively Quitting Improves LLM Agent Safety. *arXiv preprint*, 2025. arXiv:2510.16492.
- [28] Qwen Team. Qwen3.5-4B and Qwen3.5-9B (models). <https://huggingface.co/Qwen/Qwen3.5-4B>, <https://huggingface.co/Qwen/Qwen3.5-9B>.
- [29] Qwen Team. Qwen-AgentWorld-35B-A3B (model). <https://huggingface.co/Qwen/Qwen-AgentWorld-35B-A3B>.
- [30] A. Z. Ren et al. Robots That Ask For Help: Uncertainty Alignment for Large Language Model Planners. In *Proceedings of the 7th Conference on Robot Learning (CoRL)*, PMLR 229:661–682, 2023. arXiv:2307.01928.
- [31] Roundcube (software). <https://roundcube.net/>.
- [32] Y. Ruan et al. Identifying the Risks of LM Agents with an LM-Emulated Sandbox. *arXiv preprint*, 2023. arXiv:2309.15817.
- [33] C. Vicentino. Look When Unsure, Check When Sure: Consequence Training Makes a World Model’s Remaining Errors Confident, Most of All Where It Knows the World Best. *Zenodo*, 2026. DOI: 10.5281/zenodo.23146970 (all versions).
- [34] C. Vicentino. Look When Unsure, Check When Sure: Consequence Training Makes a World Model’s Remaining Errors Confident, Most of All Where It Knows the World Best. Version 2 (corrected 2026-10-06). *Zenodo*, 2026. DOI: 10.5281/zenodo.23196358.
- [35] R. Wang et al. Can Language Models Serve as Text-Based World Simulators? In *Proceedings of ACL (Short Papers)*, pp. 1–17, 2024. DOI: 10.18653/v1/2024.acl-short.1.
- [36] Z. Xiao et al. WebWorld: A Large-Scale World Model for Web Agent Training. *arXiv preprint*, 2026. arXiv:2602.14721.
- [37] E. B. Wilson. Probable Inference, the Law of Succession, and Statistical Inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927. DOI: 10.1080/01621459.1927.10502953.
- [38] O. Styles and S. Miller. WorkBench Revisited: Workplace Agents Two Years On. *arXiv preprint*, 2026. arXiv:2606.13715.
- [39] M. Wortsman et al. Robust Fine-Tuning of Zero-Shot Models. In *Proceedings of CVPR*, pp. 7949–7961, 2022. DOI: 10.1109/CVPR52688.2022.00780.
- [40] S. Yao, N. Shinn, P. Razavi, and K. Narasimhan. τ -bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains. *arXiv preprint*, 2024. arXiv:2406.12045.
- [41] X. Yu et al. SeerGuard: A Safety Framework for Mobile GUI Agents via World Model Prediction. *arXiv preprint*, 2026. arXiv:2607.15550.

- [42] K. Zhang et al. Agent Learning via Early Experience. In Proceedings of ICML, 2026. arXiv:2510.08558.
- [43] B. Zheng et al. WebGuard: Building a Generalizable Guardrail for Web Agents. arXiv preprint, 2025. arXiv:2507.14293.
- [44] Y. Zuo et al. Qwen-AgentWorld: Language World Models for General Agents. arXiv preprint, 2026. arXiv:2606.24597.